

Assignment Thirteen

CS 499

Richard McCormick (RLM443)

Python Program:

```
# <-- BEGIN IMPORTS / HEADERS -->
import os
import urllib
import urllib.request
import pandas as pd
import numpy as np
import plotnine as p9
import torch
import torchvision

import sklearn
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split

from statistics import mode
import inspect
import warnings

import math
# <-- END IMPORTS / HEADERS -->

# <-- BEGIN INITIALIZATION -->
# FILE VARIABLES
download_directory = "."

# - Zip data (Training) variables
ziptrain_url = "https://hastie.su.domains/ElemStatLearn/datasets/zip.train.gz"
ziptrain_file = "zip.train.gz"
ziptrain_file_path = os.path.join(download_directory, ziptrain_file)

# - Zip data (Test) variables
ziptest_url = "https://hastie.su.domains/ElemStatLearn/datasets/zip.test.gz"
```

```

ziptest_file = "zip.test.gz"
ziptest_file_path = os.path.join(download_directory, ziptest_file)

# CONSTANT VARIABLES
spam_label_col = 57
zip_empty_col = 257

MAX_EPOCHS_VAR = 10000
BATCH_SIZE_VAR = 512
STEP_SIZE_VAR = 0.01
HIDDEN_LAYERS_VAR = 1000000000000000
CV_VAL = 2
N_FOLDS = 2

device = "cuda" if torch.cuda.is_available() else "cpu"

global ncol
global n_classes

# MISC. VARIABLES
kf = KFold( n_splits=N_FOLDS, shuffle=True, random_state=1 )
test_acc_df_list = []
pipe = make_pipeline(StandardScaler(), LogisticRegression(max_iter=1000))

#CLASS DEFINITIONS
class CSV(torch.utils.data.Dataset):
    def __init__(self, features, labels):
        self.features = features
        self.labels = labels
    def __getitem__(self, item):
        return self.features[item,:], self.labels[item]
    def __len__(self):
        return len(self.labels)

class TorchConvModel(torch.nn.Module):
    def __init__(self, *units_per_layer):
        super(TorchConvModel, self).__init__()
        seq_args = []

        seq_args.append( torch.nn.Conv2d( in_channels=1,
                                           out_channels=32,
                                           kernel_size=3,
                                           stride=3 ) )

        seq_args.append( torch.nn.ReLU() )

```

```

seq_args.append( torch.nn.Conv2d( in_channels=32,
                                   out_channels=64,
                                   kernel_size=3,
                                   stride=3 ) )

seq_args.append( torch.nn.ReLU() )

seq_args.append( torch.nn.MaxPool2d(kernel_size=1, stride=1) )

seq_args.append( torch.nn.ReLU() )

seq_args.append( torch.nn.Flatten(start_dim=1) )

seq_args.append( torch.nn.Linear( 64, 128 ) )

seq_args.append( torch.nn.ReLU() )

seq_args.append( torch.nn.Linear( 128, 2 ) )

self.stack = torch.nn.Sequential(*seq_args)

def forward(self, feature_mat):
    return self.stack(feature_mat.float())

class TorchModel(torch.nn.Module):
    def __init__(self, *units_per_layer):
        super(TorchModel, self).__init__()
        seq_args = []
        for layer_i in range(len(units_per_layer)-1):
            units_in = units_per_layer[layer_i]
            units_out = units_per_layer[layer_i+1]
            seq_args.append( torch.nn.Linear( units_in, units_out ) )
            if layer_i != len(units_per_layer)-2:
                seq_args.append(torch.nn.ReLU())
        self.stack = torch.nn.Sequential(*seq_args)

    def forward(self, feature_mat):
        return self.stack(feature_mat.float())

class ConvolutionalMPL:
    def __init__(self, **kwargs):
        """Store hyper-parameters, TorchConvModel instance, loss, etc."""
        kwargs.setdefault("max_epochs", 2)
        kwargs.setdefault("batch_size", BATCH_SIZE_VAR)

```

```

kwargs.setdefault("step_size", 0.01)
kwargs.setdefault("units_per_layer", ( ncol, 1000, 100, n_classes ) )
kwargs.setdefault("hidden_layers", 3)
kwargs.setdefault("opt_name", torch.optim.SGD)
kwargs.setdefault("opt_params", {'lr':0.1})

for key, value in kwargs.items():
    setattr(self, key, value)

units_per_layer = [ncol]
for L in range(self.hidden_layers):
    units_per_layer.append(1024)
units_per_layer.append(n_classes)

self.best_epoch = -1                # Best Epoch
self.loss_df = pd.DataFrame()       # Dataframe of Loss per Epoch

self.model = TorchConvModel(*self.units_per_layer).to(device)

self.optimizer = self.opt_name(self.model.parameters(), **self.opt_params)
self.loss_fun = torch.nn.CrossEntropyLoss()

def take_step(self, X, y):
    """compute predictions, loss, gradients, take one step"""
    self.optimizer.zero_grad()
    pred_tensor = self.model.forward(X).reshape(len(y))
    loss_tensor = self.loss_fun(pred_tensor, y.long())
    loss_tensor.backward()
    self.optimizer.step()

def fit(self, X, y):
    """Gradient descent learning of weights"""
    units_per_layer = [ncol]
    for L in range(self.hidden_layers):
        units_per_layer.append(100)
    units_per_layer.append(n_classes)

    feature_tensor = X.reshape(len(X),1,16,16)

    ds = CSV( feature_tensor, y )
    dl = torch.utils.data.DataLoader( ds, batch_size = self.batch_size,
                                      shuffle = True )

    loss_df_list = []
    best_loss_val = 10000

```

```

for epoch in range(self.max_epochs):
    for batch_features, batch_labels in dl:
        self.take_step(batch_features, batch_labels)
        pred = self.model(batch_features)
        loss_value = self.loss_fun(pred, batch_labels.long())

        if( loss_value < best_loss_val ):
            self.best_epoch = epoch
            best_loss_val = loss_value

    loss_df_list.append(pd.DataFrame({
        #"set_name":set_name,
        "loss":float(loss_value),
        "epoch":epoch,
    }, index=[0]))#subtrain/validation loss using current weights.

self.loss_df = pd.concat( loss_df_list )

def predict(self, X):
    """Return numpy vector of predictions"""
    pred_vec = []
    for row in self.model(torch.from_numpy(X.reshape(len(X),1,16,16))):
        best_label = -1
        highest_prob = -1000
        itera = 0
        for iter in row.long():
            if(iter.item() > highest_prob):
                highest_prob = iter.item()
                best_label = itera
            itera += 1
        pred_vec.append(best_label)

    return pred_vec

class TorchConvLearnerCV:
    def __init__(self, max_epochs, batch_size, step_size, units_per_layer,
**kwargs):
        self.subtrain_learner = ConvolutionalMLP( max_epochs=max_epochs,
                                                    batch_size=batch_size,
                                                    step_size=step_size,
                                                    units_per_layer=units_per_layer )

        for key, value in kwargs.items():
            setattr(self, key, value)

```

```

self.batch_size = batch_size
self.step_size = step_size
self.units_per_layer = units_per_layer

self.plotting_df = pd.DataFrame()

def fit(self, X, y):
    """cross-validation for selecting the best number of epochs"""
    fold_vec = np.random.randint(low=0, high=5, size=y.size)
    validation_fold = 0
    is_set_dict = {
        "validation": fold_vec == validation_fold,
        "subtrain": fold_vec != validation_fold,
    }

    set_features = {}
    set_labels = {}

    for set_name, is_set in is_set_dict.items():
        set_features[set_name] = X[is_set, :]
        set_labels[set_name] = y[is_set]
    {set_name: array.shape for set_name, array in set_features.items()}

    self.subtrain_learner.validation_data = set_features["validation"]
    self.subtrain_learner.fit( set_features["subtrain"],
set_labels["subtrain"], "subtrain" )
    self.plotting_df = pd.concat([self.plotting_df,
self.subtrain_learner.loss_df])

    best_epochs = self.subtrain_learner.best_epoch

    self.train_learner = ConvolutionalMLP( max_epochs=best_epochs,
batch_size=self.batch_size,
step_size=self.step_size,
units_per_layer=self.units_per_layer )
    self.train_learner.fit( set_features["validation"],
set_labels["validation"], "validation" )
    self.plotting_df = pd.concat([self.plotting_df,
self.train_learner.loss_df])

    def predict(self, X):
        return self.train_learner.predict(X)

class RegularizedMLP:
    def __init__(self, **kwargs):

```

```

    kwargs.setdefault("max_epochs", MAX_EPOCHS_VAR)
    kwargs.setdefault("batch_size", BATCH_SIZE_VAR)
    kwargs.setdefault("step_size", 0.001)
    kwargs.setdefault("units_per_layer", ( ncol, 1000, 100, n_classes ) )

    for key, value in kwargs.items():
        setattr(self, key, value)

    self.estimator = estimator()

def fit(self, X, y):
    self.estimator.fit(X, y)

def predict(self, X):
    return self.estimator.predict(X)

class MyCV():
    def __init__(self, **kwargs):
        # Initialize parameters and setup variables
        self.train_features = []
        self.train_labels = []
        self.training_data = None

        kwargs.setdefault("num_folds", 5)
        for key, value in kwargs.items():
            setattr(self, key, value)

        self.estimator = self.estimator()
        self.best_model = None

        self.plotting_df = pd.DataFrame()

    def fit(self, X, y):
        # Populate internal data structures
        self.train_features = X
        self.train_labels = y
        self.training_data = {'X':self.train_features, 'y':self.train_labels}

        # Create a dataframe to temporarily hold results from each fold
        best_paramter_df = pd.DataFrame()

        # Calculate folds
        fold_indicies = []

        # Pick random entries for validation/subtrain

```

```

fold_vec = np.random.randint(low=0,
                             high=self.num_folds,
                             size=self.train_labels.size)

# for each fold,
for fold_number in range(self.num_folds):
    subtrain_indicies = []
    validation_indicies = []
    # check if index goes into subtrain or validation list
    for index in range(len(self.train_features)):
        if fold_vec[index] == fold_number:
            validation_indicies.append(index)
        else:
            subtrain_indicies.append(index)

    fold_indicies.append([subtrain_indicies, validation_indicies])

printing_df = pd.DataFrame()
# Loop over the folds
for foldnum, indicies in enumerate(fold_indicies):
    #print("(MyCV) Subfold #" + str(foldnum))

    # Get indicies of data chosen for this fold
    index_dict = dict(zip(["subtrain", "validation"], indicies))
    set_data_dict = {}

    # Dictionary for test and train data
    for set_name, index_vec in index_dict.items():
        set_data_dict[set_name] = {
            "X":self.train_features[index_vec],
            "y":self.train_labels[index_vec]
        }

    # Create a dictionary to hold the results of the fitting
    results_dict = {}

    parameter_index = 0
    # Loop over each parameter in the param_grid
    for parameter_entry in self.param_grid:
        for param_name, param_value in parameter_entry.items():
            setattr(self.estimator, param_name, param_value)

            # Fit fold data to estimator
            self.estimator.fit(**set_data_dict["subtrain"])

```



```

        printing_df = self.estimator.loss_df
        for param_name, param_value in parameter_entry.items():
            printing_df[param_name] = str(param_value)
        printing_df['set'] = 'subtrain'
        printing_df['subfold'] = foldnum
        self.plotting_df = pd.concat([self.plotting_df, printing_df])

        self.estimator.fit(**set_data_dict["validation"])

        printing_df = self.estimator.loss_df
        for param_name, param_value in parameter_entry.items():
            printing_df[param_name] = str(param_value)
        printing_df['set'] = 'validation'
        printing_df['subfold'] = foldnum
        self.plotting_df = pd.concat([self.plotting_df, printing_df])

        # Make a prediction of current fold's test data
        prediction = \
            self.estimator.predict(set_data_dict["validation"]['X'])

        # Determine accuracy of the prediction
        results_dict[parameter_index] = \
            (prediction == set_data_dict["validation"]['y']).mean()*100

        # index only serves to act as key for results dictionary
        parameter_index += 1

    # Store the results of this param entry into dataframe
    best_paramter_df = best_paramter_df.append(results_dict,
                                                ignore_index=True)

# Average across all folds for each parameter
averaged_results = dict(best_paramter_df.mean())

# From the averaged data, get the single best model
best_result = max(averaged_results, key = averaged_results.get)

# Store best model for future reference
self.best_model = self.param_grid[best_result]

def predict(self, test_features):
    # Load best model into estimator
    for param_name, param_value in self.best_model.items():
        setattr(self.estimator, param_name, param_value)

```

```

        # Fit estimator to training data
        self.estimator.fit(**self.training_data)

        # Make a prediction of the test features
        prediction = self.estimator.predict(test_features)

        return(prediction)

class OptimizerMLP:
    def __init__(self, **kwargs):
        """Store hyper-parameters, TorchModel instance, loss, etc."""
        kwargs.setdefault("max_epochs", 2)
        kwargs.setdefault("batch_size", BATCH_SIZE_VAR)
        kwargs.setdefault("step_size", 0.01)
        kwargs.setdefault("units_per_layer", ( ncol, 95, 10, 128, 2 ) )
        kwargs.setdefault("hidden_layers", 3)
        kwargs.setdefault("opt_name", torch.optim.SGD)
        kwargs.setdefault("opt_params", {'lr':0.1})

        for key, value in kwargs.items():
            setattr(self, key, value)

        units_per_layer = [ncol]
        for L in range(self.hidden_layers):
            units_per_layer.append(100)
        units_per_layer.append(n_classes)

        self.best_epoch = -1                # Best Epoch
        self.loss_df = pd.DataFrame()       # Dataframe of Loss per Epoch

        self.model = TorchModel(*self.units_per_layer).to(device)

        self.optimizer = self.opt_name(self.model.parameters(), **self.opt_params)
        self.loss_fun = torch.nn.CrossEntropyLoss()

    def take_step(self, X, y):
        """compute predictions, loss, gradients, take one step"""
        self.optimizer.zero_grad()
        pred_tensor = self.model.forward(X).reshape(len(y))
        loss_tensor = self.loss_fun(pred_tensor, y.long())
        loss_tensor.backward()
        self.optimizer.step()

    def fit(self, X, y):

```

```

"""Gradient descent learning of weights"""
units_per_layer = [ncol]
for L in range(self.hidden_layers):
    units_per_layer.append(100)
units_per_layer.append(n_classes)

ds = CSV( X, y )
dl = torch.utils.data.DataLoader( ds, batch_size = self.batch_size,
                                shuffle = True )

loss_df_list = []
best_loss_val = 10000

for epoch in range(self.max_epochs):
    for batch_features, batch_labels in dl:
        self.take_step(batch_features, batch_labels)
        pred = self.model(batch_features)
        loss_value = self.loss_fun(pred, batch_labels.long())

        if( loss_value < best_loss_val ):
            self.best_epoch = epoch
            best_loss_val = loss_value

    loss_df_list.append(pd.DataFrame({
        "set_name":set_name,
        "loss":float(loss_value),
        "epoch":epoch,
    }, index=[0]))#subtrain/validation loss using current weights.

self.loss_df = pd.concat( loss_df_list )

def predict(self, X):
    """Return numpy vector of predictions"""
    pred_vec = []
    for row in self.model(torch.from_numpy(X)):
        best_label = -1
        highest_prob = -1000
        itera = 0
        for iter in row.long():
            if(iter.item() > highest_prob):
                highest_prob = iter.item()
                best_label = itera
            itera += 1
        pred_vec.append(best_label)

    return pred_vec

```

[illegible]

```

                                step_size=self.step_size,
                                units_per_layer=self.units_per_layer )
        self.train_learner.fit( set_features["validation"],
set_labels["validation"], "validation" )
        self.plotting_df = pd.concat([self.plotting_df,
self.train_learner.loss_df])

    def predict(self, X):
        return self.train_learner.predict(X)

# <-- END INITIALIZATION -->

# <-- BEGIN FUNCTIONS -->
# FUNCTION: MAIN
# Description : Main driver for Assignment Ten
# Inputs      : None
# Outputs     : PlotNine graphs, printed and saved to directory
# Dependencies : build_image_df_from_dataframe
def main():
    # Display the title
    print("\nCS 499: Homework 14 Program Start")
    print("=====\n")

    # Monsoon stuff
    none_or_str = os.getenv("SLURM_JOB_CPUS_PER_NODE")
    CPUS = int(1 if none_or_str is None else none_or_str)
    torch.set_num_interop_threads(CPUS)
    torch.set_num_threads(CPUS)

    # Suppress annoying plotnine warnings
    warnings.filterwarnings('ignore')

    # Download data files
    download_data_file(ziptrain_file, ziptrain_url, ziptrain_file_path)
    download_data_file(ziptest_file, ziptest_url, ziptest_file_path)

    # Open each dataset as a pandas dataframe
    zip_train_df = pd.read_csv(ziptrain_file, header=None, sep=" ")
    zip_test_df = pd.read_csv(ziptest_file, header=None, sep=" ")

    # Concat the two zip dataframes together
    zip_df = pd.concat([zip_train_df, zip_test_df])
    zip_df[0] = zip_df[0].astype(int)

    # Drop empty col from zip dataframe

```

```

zip_df = zip_df.drop(columns=[zip_empty_col])

zip_ones = zip_df[zip_df[0].isin([0, 1])]
zip_sevens = zip_df[zip_df[0].isin([0, 7])]

zip_one_features = zip_ones.iloc[:, :-1].to_numpy()
zip_one_labels = zip_ones[0].to_numpy()

zip_seven_features = zip_sevens.iloc[:, :-1].to_numpy()
zip_seven_labels = zip_sevens[0].to_numpy()

zip_seven_labels = np.where(zip_sevens[0] == 7, 1, 0)

# Create data dictionary
data_dict = {
    'zip_seven_zero' : [zip_seven_features, zip_seven_labels],
    'zip_one_zero' : [zip_one_features, zip_one_labels]
}

final_df_list = []
final_deep_print_list = []

final_deep_df = pd.DataFrame()

# Loop through each data set
for data_set, (input_data, output_array) in data_dict.items():
    current_set = str(data_set)
    #print("")
    #print("Working on set: " + current_set)

    # Loop over each fold for each data set
    for foldnum, indices in enumerate(kf.split(input_data)):
        #print("Fold #" + str(foldnum))

        # Set up input data structs
        global ncol
        nrow, ncol = input_data.shape
        index_dict = dict(zip(["train", "test"], indices))

        # Creating dictionary with input and outputs
        set_data_dict = {}
        for set_name, index_vec in index_dict.items():
            set_data_dict[set_name] = {
                "X":input_data[index_vec],
                "y":output_array[index_vec]
            }

```

```

    }

    # Finalizing variables for CV construction
    param_dicts = [{'n_neighbors':[x]} for x in range(1, 21)]
    global n_classes
    n_classes = len( np.unique( set_data_dict['test']['y'] ) )
    UNITS_PER_VAR = ( 256, 512, 1024, 512, 1024, 512, 2048, 128, 1 )

    param_grid = []

    for momentum in 0.1, 0.5, 0.9:
        for lr in 0.1, 0.01, 0.001:
            param_grid.append({
                "opt_name":torch.optim.SGD,
                "opt_params":{"momentum":momentum, "lr":lr}
            })
    for beta1 in 0.85, 0.9, 0.95:
        for beta2 in 0.99, 0.999, 0.9999:
            param_grid.append({
                "opt_name":torch.optim.Adam,
                "opt_params":{"betas":(beta1, beta2)}
            })

    clf = GridSearchCV(KNeighborsClassifier(), param_dicts)
    linear_model = sklearn.linear_model.LogisticRegressionCV(cv=CV_VAL)
    DeepConvTorchCV = MyCV( max_epochs = MAX_EPOCHS_VAR,
                            batch_size = BATCH_SIZE_VAR,
                            step_size = STEP_SIZE_VAR,
                            units_per_layer = UNITS_PER_VAR,
                            estimator = ConvolutionalMLP,
                            param_grid = param_grid,
                            num_folds = CV_VAL )

    DeepTorchCV = MyCV( max_epochs = MAX_EPOCHS_VAR,
                        batch_size = BATCH_SIZE_VAR,
                        step_size = STEP_SIZE_VAR,
                        units_per_layer = UNITS_PER_VAR,
                        estimator = OptimizerMLP,
                        param_grid = param_grid,
                        num_folds = CV_VAL )

    # Train the models with given data
    clf.fit(**set_data_dict["train"])
    linear_model.fit(**set_data_dict["train"])
    DeepConvTorchCV.fit(**set_data_dict["train"])

```

```

DeepTorchCV.fit(**set_data_dict["train"])

# Get most common output from outputs for featureless set
most_common_element = mode(set_data_dict["train"]['y'])

buffer_df = DeepConvTorchCV.plotting_df
buffer_df['fold'] = foldnum
buffer_df['data_set'] = data_set
final_deep_print_list.append(buffer_df)

# Get results
pred_dict = {
    "GridSearchCV + KNeighborsClassifier": \
        clf.predict(set_data_dict["test"]["X"]),
    "LogisticRegressionCV": \
        linear_model.predict(set_data_dict["test"]["X"]),
    "ConvolutionalMLP": \
        DeepConvTorchCV.predict(set_data_dict["test"]["X"]),
    "DenseMLP": \
        DeepTorchCV.predict(set_data_dict["test"]["X"]),
    "Featureless":most_common_element
}

# Build results dataframe for each algo/fold
for algorithm, pred_vec in pred_dict.items():
    test_acc_dict = {
        "test_accuracy_percent":(
            pred_vec == set_data_dict["test"]["y"]).mean()*100,
        "data_set":data_set,
        "fold_id":foldnum,
        "algorithm":algorithm
    }
    test_acc_df_list.append(pd.DataFrame(test_acc_dict, index=[0]))

final_deep_df = pd.concat(final_deep_print_list)

# Build accuracy results dataframe
test_acc_df = pd.concat(test_acc_df_list)

# Print results
print("\n")
print(test_acc_df)
print("")

# Plot results

```



```

plot = (p9.ggplot(test_acc_df,
                 p9.aes(x='test_accuracy_percent',
                       y='algorithm'))
      + p9.facet_grid('. ~ data_set')
      + p9.geom_point()
      + p9.theme(subplots_adjust={'left': 0.2}))

# Epoch vector for plotting
"""epoch_vec = np.arange(MAX_EPOCHS_VAR)
epoch_vec = np.tile(epoch_vec, 1)
epoch_vec = epoch_vec.flatten()"""

final_deep_df['opt_and_params'] = final_deep_df['opt_name'] +
final_deep_df['opt_params']
final_deep_df = final_deep_df.groupby(['set', 'data_set', 'epoch',
'opt_and_params'], as_index=False).mean()
#final_deep_df['epochs'] = epoch_vec
deepplot = (p9.ggplot(final_deep_df,
                    p9.aes(x='epoch',
                          y='loss',
                          color='set'))
          + p9.facet_grid('opt_and_params ~ data_set', scales='free',
shrink=True)
          + p9.geom_line()
          + p9.theme(subplots_adjust={'left': 0.2},
                    strip_text_y = p9.element_text(angle = 0,ha =
'left'),
                    figure_size=(25,25))
          + p9.ggtitle("Loss vs. Epochs"))

plot.save("accuracy_plot.png")
deepplot.save("DeepTorch Loss Graph.png")

print("\nCS 499: Homework 14 Program End")
print("=====\n")

# FUNCTION : DOWNLOAD_DATA_FILE
# Description: Downloads file from source, if not already downloaded
# Inputs:
#   - file      : Name of file to download
#   - file_url  : URL of file
#   - file_path : Absolute path of location to download file to.
#                 Defaults to the local directory of this program.
# Outputs: None
def download_data_file(file, file_url, file_path):

```

```

# Check for data file. If not found, download
if not os.path.isfile(file_path):
    try:
        print("Getting file: " + str(file) + "...\\n")
        urllib.request.urlretrieve(file_url, file_path)
        print("File downloaded.\\n")
    except(error):
        print(error)
else:
    print("File: " + str(file) + " is already downloaded.\\n")

# FUNCTION : DOWNLOAD_DATA_FILE
# Description: Downloads file from source, if not already downloaded
# Inputs:
#     - file      : Name of file to download
# Outputs: None
def get_n_params(module):
    return sum(
        [math.prod(list(p.shape)) for p in module.parameters()])

# Launch main
if __name__ == "__main__":
    main()

```

Program Output:

```
(cs499f22) [rlm443@wind ~]$ time srun -t 1:00:00 --mem=8gb --cpus-per-task=1 python Homework14.py

CS 499: Homework 13 Program Start
=====

File: zip.train.gz is already downloaded.

File: zip.test.gz is already downloaded.


  test_accuracy_percent  ...      algorithm
0          100.000000  ...  GridSearchCV + KNeighborsClassifier
0          100.000000  ...      LogisticRegressionCV
0          100.000000  ...      ConvolutionalMLP
0          100.000000  ...      DenseMLP
0          66.410912  ...      Featureless
0          100.000000  ...  GridSearchCV + KNeighborsClassifier
0          100.000000  ...      LogisticRegressionCV
0          100.000000  ...      ConvolutionalMLP
0          100.000000  ...      DenseMLP
0          66.040956  ...      Featureless
0          99.858257  ...  GridSearchCV + KNeighborsClassifier
0          99.858257  ...      LogisticRegressionCV
0          99.645641  ...      ConvolutionalMLP
0          99.787385  ...      DenseMLP
0          55.067328  ...      Featureless
0          100.000000  ...  GridSearchCV + KNeighborsClassifier
0          99.858257  ...      LogisticRegressionCV
0          99.291283  ...      ConvolutionalMLP
0          99.787385  ...      DenseMLP
0          54.996456  ...      Featureless

[20 rows x 4 columns]

CS 499: Homework 13 Program End
=====

real    0m44.899s
user    0m0.010s
sys     0m0.011s
```

```

(cs499f22) [rlm443@wind ~ ]$ time srun -t 1:00:00 --gres=gpu:tesla:1 --mem=8gb --cpus-per-task=1 python Homework14.py

CS 499: Homework 13 Program Start
=====

File: zip.train.gz is already downloaded.

File: zip.test.gz is already downloaded.


    test_accuracy_percent ...      algorithm
0      100.000000 ...  GridSearchCV + KNeighborsClassifier
0      100.000000 ...      LogisticRegressionCV
0      100.000000 ...      ConvolutionalMLP
0      100.000000 ...      DenseMLP
0      66.410912 ...      Featureless
0      100.000000 ...  GridSearchCV + KNeighborsClassifier
0      100.000000 ...      LogisticRegressionCV
0      100.000000 ...      ConvolutionalMLP
0      100.000000 ...      DenseMLP
0      66.040956 ...      Featureless
0      99.858257 ...  GridSearchCV + KNeighborsClassifier
0      99.858257 ...      LogisticRegressionCV
0      99.220411 ...      ConvolutionalMLP
0      99.858257 ...      DenseMLP
0      55.067328 ...      Featureless
0      100.000000 ...  GridSearchCV + KNeighborsClassifier
0      99.858257 ...      LogisticRegressionCV
0      99.433026 ...      ConvolutionalMLP
0      99.787385 ...      DenseMLP
0      54.996456 ...      Featureless

[20 rows x 4 columns]

CS 499: Homework 13 Program End
=====

real    0m38.443s
user    0m0.012s
sys     0m0.006s

```

```
(cs499f22) [rlm443@wind ~]$ time srun -t 1:00:00 --mem=8gb --cpus-per-task=4 python Homework14.py
```

```
CS 499: Homework 13 Program Start
```

```
=====
```

```
File: zip.train.gz is already downloaded.
```

```
File: zip.test.gz is already downloaded.
```

	test_accuracy_percent	...	algorithm
0	100.000000	...	GridSearchCV + KNeighborsClassifier
0	100.000000	...	LogisticRegressionCV
0	100.000000	...	ConvolutionalMLP
0	100.000000	...	DenseMLP
0	66.410912	...	Featureless
0	100.000000	...	GridSearchCV + KNeighborsClassifier
0	100.000000	...	LogisticRegressionCV
0	100.000000	...	ConvolutionalMLP
0	100.000000	...	DenseMLP
0	66.040956	...	Featureless
0	99.858257	...	GridSearchCV + KNeighborsClassifier
0	99.858257	...	LogisticRegressionCV
0	99.574770	...	ConvolutionalMLP
0	99.787385	...	DenseMLP
0	55.067328	...	Featureless
0	100.000000	...	GridSearchCV + KNeighborsClassifier
0	99.858257	...	LogisticRegressionCV
0	99.645641	...	ConvolutionalMLP
0	99.787385	...	DenseMLP
0	54.996456	...	Featureless

```
[20 rows x 4 columns]
```

```
CS 499: Homework 13 Program End
```

```
=====
```

```
real    0m36.008s
user    0m0.012s
sys     0m0.007s
```

Question Answers / Commentary:

For this assignment I was able to modify my Homework 13 assignment to be able to run on Monsoon. Running the program on a CPU and a GPU made a significant difference in the time it took for the program to finish. Running the program on a GPU showed around a 15% speedup over running the same program on a CPU. Likewise, running the program on a CPU with 4 cores instead of 1 saw a speedup of around 20%.

I found that there was a sweet spot in the number of cores I could allocate to the program and still see a speedup. For example, running the program with 4 cores saw a significant speedup, but running the same program with 8 cores caused it to take longer. I believe that this is due to my program not having multithreading options properly configured.

The test accuracy for all three of the above programs is roughly the same. Between them there are variations of around 0.5% in accuracy. This could be due to the different cores running processes in slightly different ways, or due to slight changes in how the Cross Validation folds were done. Regardless, the program seems to maintain its original high accuracy across the board.